



# Mongo Database



*Scale the web*


# Course Objectives

---



- Learn how to use Mongo database with the main operations and how to deal with the mongo shell.

# Course Evaluation



# Course Prerequisites

---



- Database Fundamentals
- Familiarity with Linux terminal
- Basic Knowledge of Javascript

# Course Material

---



You can access course material via this URL:

<http://tinyurl.com/iti-mongodb>


# Agenda

---



- What is NoSQL ?
- Why NoSQL?
- Why MongoDB ?
- Getting Started
- Mongo Basics

# What is NoSQL



# What is NoSQL



## Definitions:

- **NoSQL** is a movement promoting a loosely defined class of non-relational data stores that break with a long history of relational databases.
- Non-relational next generation operational data-stores and databases.
- Next Generation Databases mostly addressing some of the points: being non-relational, distributed, open-source and horizontal scalable.
- NoSQL is a term coined by **Carlo Strozzi** and repurposed by **Eric Evans** to refer to “some” storage systems.

The term should be used as  
**NOT ONLY SQL**



# What is NoSQL



## Types:

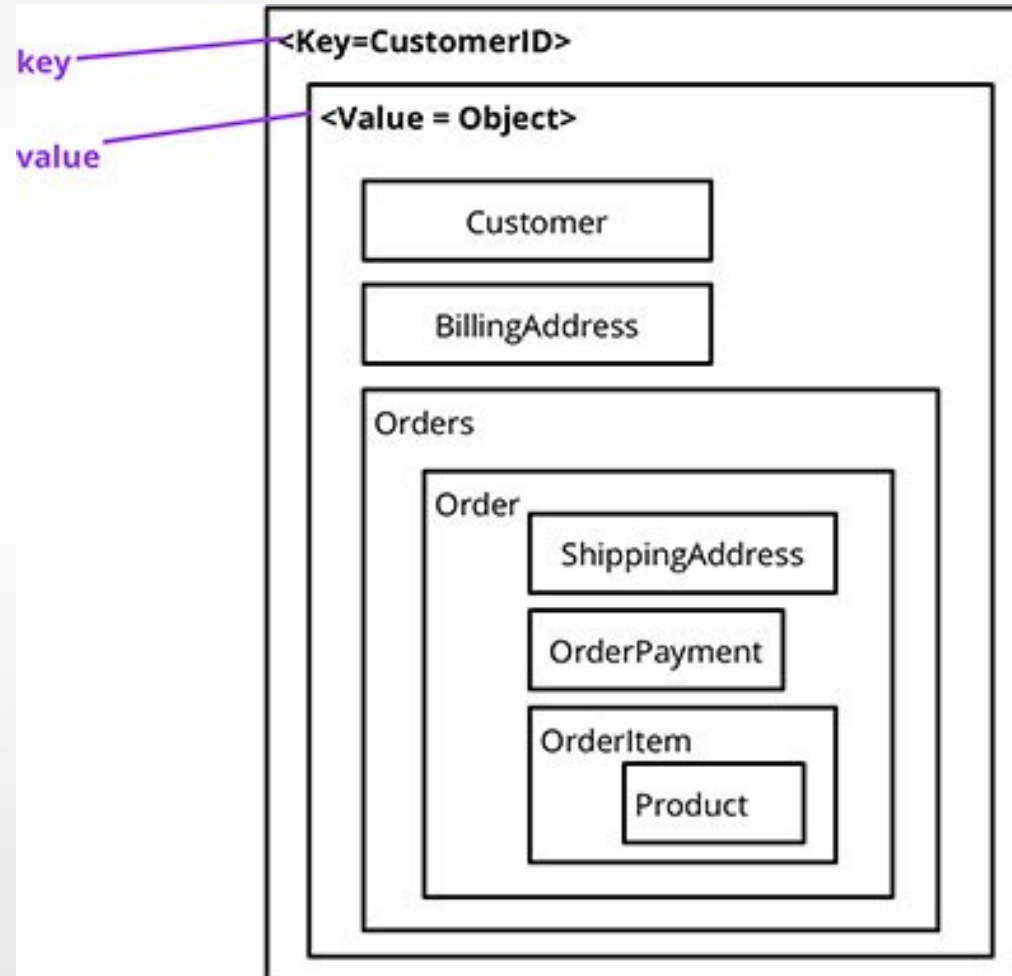
There have been various approaches to classify NoSQL databases

- **Column:** Cassandra
- **Document:** MongoDB
- **Key-value:** FoundationDB
- **Graph:** Neo4J

# What is NoSQL



## Types:



# What is NoSQL



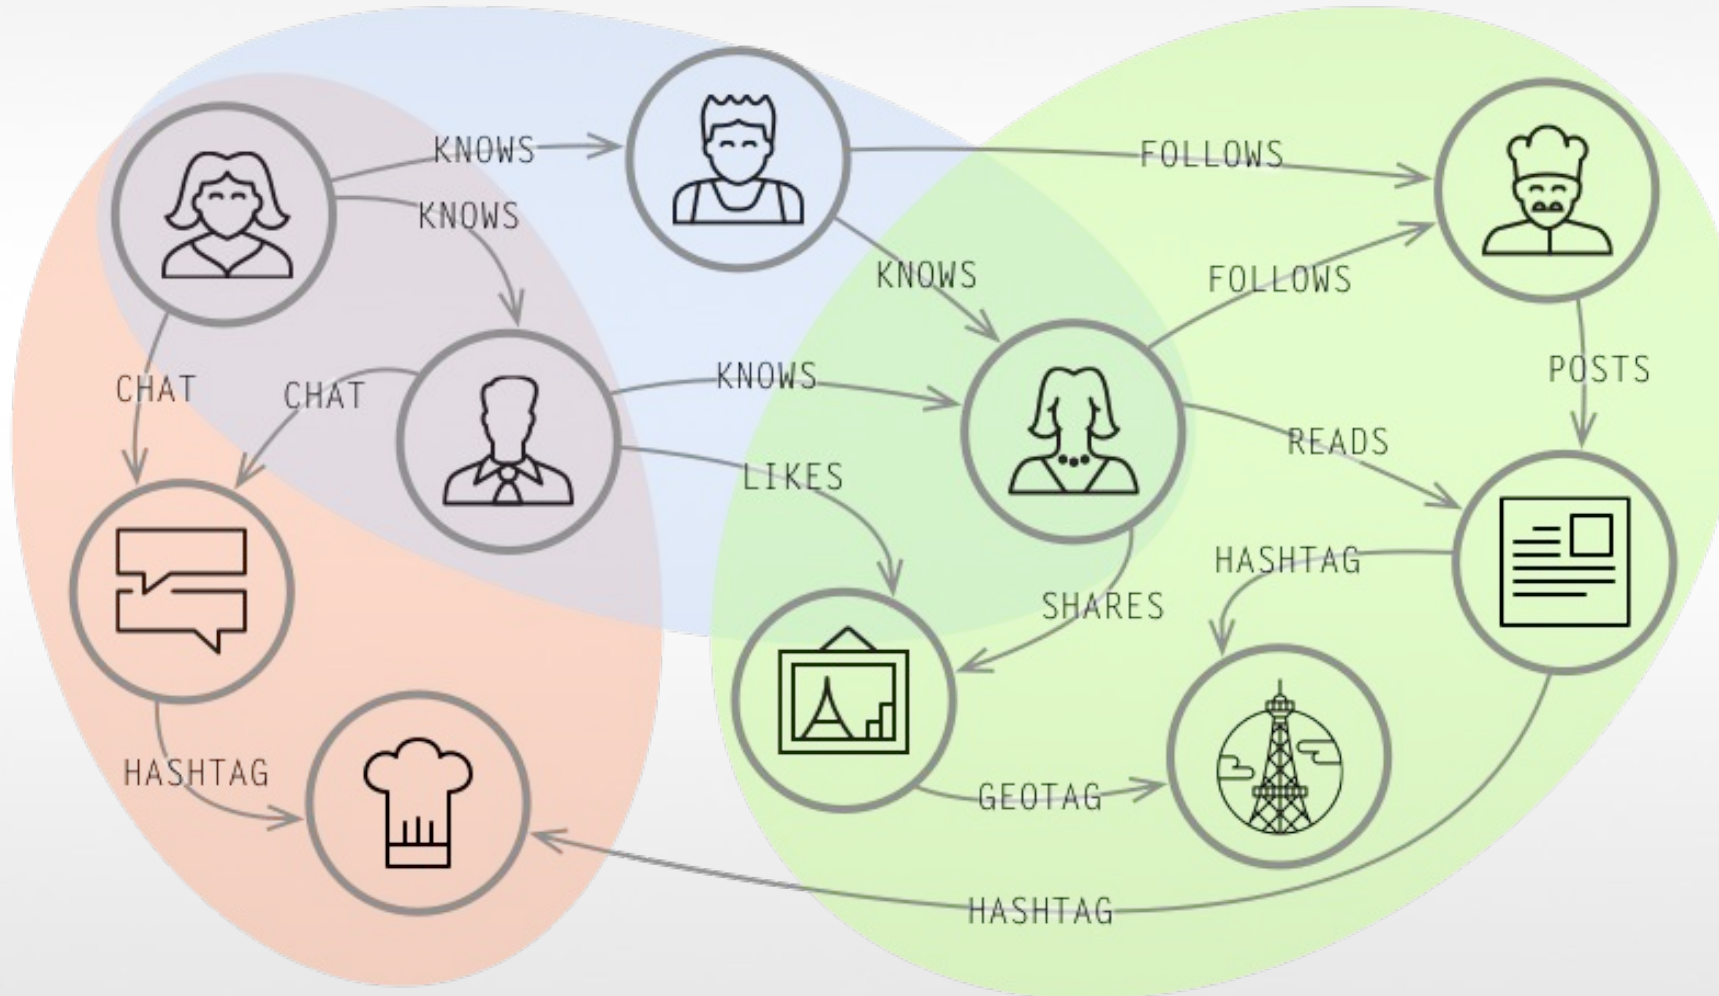
## Types:



# What is NoSQL



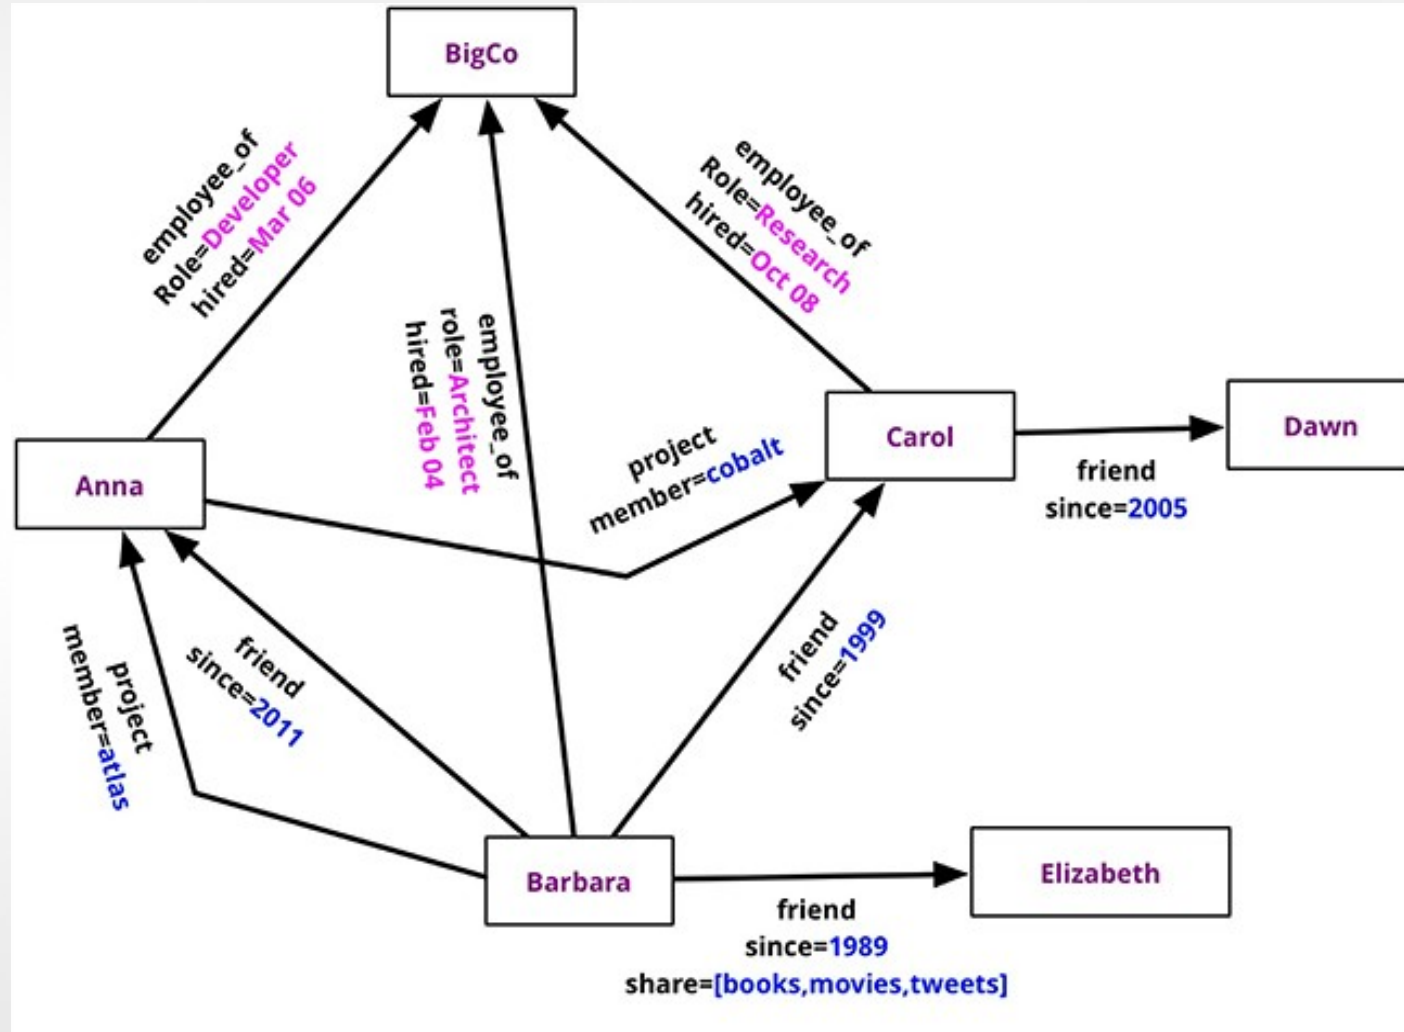
## Types:



# What is NoSQL



Types:





# Why NoSQL?

# Why NoSQL



- **Big Data** is one of the key forces driving the growth and popularity of **NoSQL** for business.
- **A Big Data project is normally typified by:**
  - High data velocity:** lots of data coming in very quickly, possibly from different locations.
  - Data variety:** storage of data that is structured, semi-structured and unstructured.
  - Data volume:** data that involves many terabytes or petabytes in size.
  - Data complexity:** data that is stored and managed in different locations or data centers.


# Why NoSQL



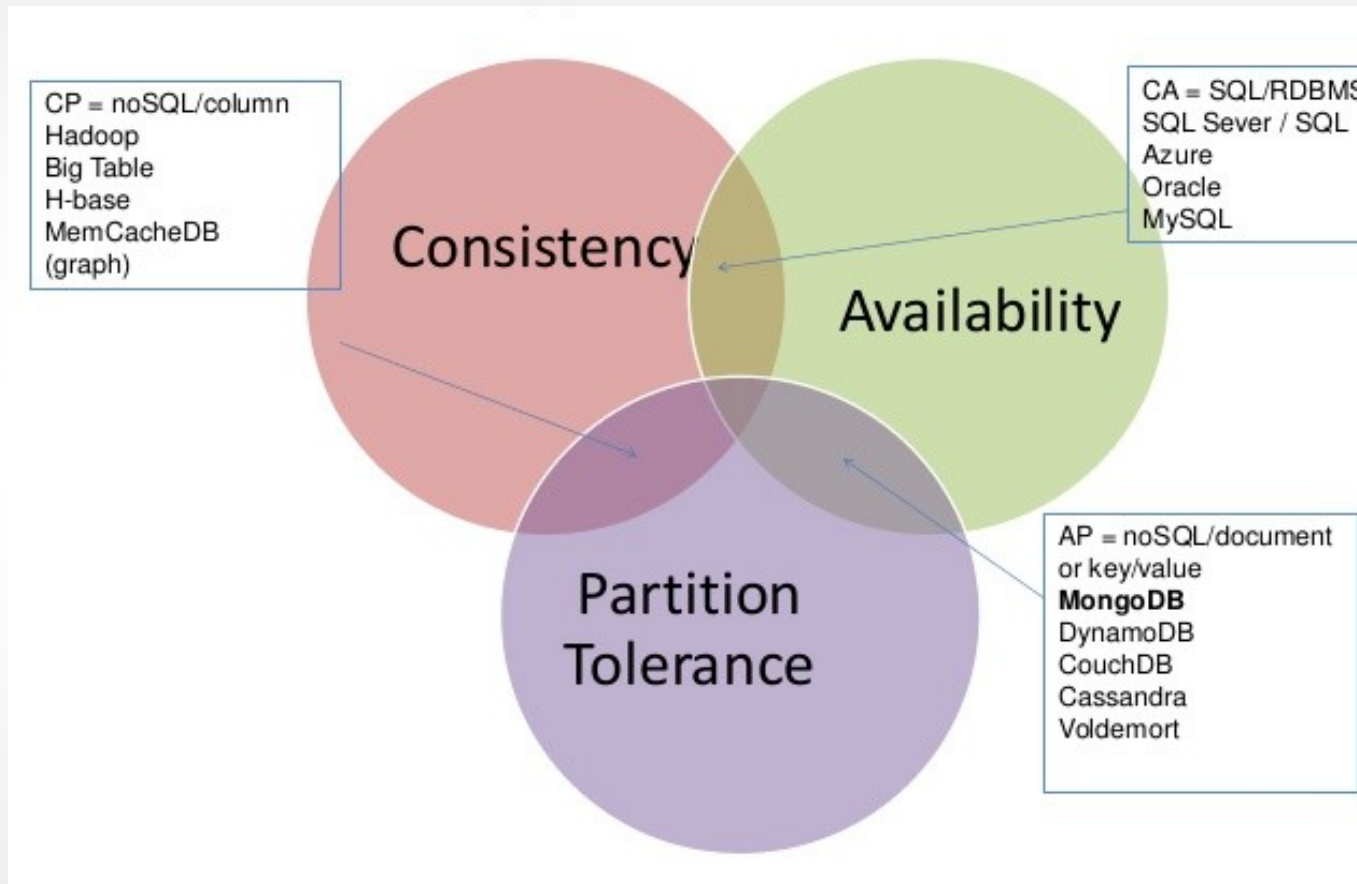

**That is why we need NoSQL**



# Why NoSQL



# Why NoSQL



## CAP Theory

# Why NoSQL



## Scenarios where NoSQL **SHOULD** be used:

- Your relational database will not scale to your traffic at an acceptable cost.
- In a NoSQL database, there is no fixed schema and no joins. NoSQL can take advantage of “scaling out”. Scaling out refers to spreading the load over many commodity systems.
- It’s useful for creating prototypes or fast applications as it provides a tool to develop new features easily.
- You have local data transactions which do not have to be very durable. **e.g. “liking”** items on websites.
- Agile sprints, quick iteration, and frequent code pushes
- Object-oriented programming that is easy to use and flexible

# Why NoSQL



## Scenarios where NoSQL **SHOULD NOT** be used:

- It cannot necessarily guarantee the ACID(Atomicity,Consistency,Isolation,Durability) properties for your transactions.
- Normally an interface is provided for storing your data. Do not try to use a complicated query in that interface.
- The developer should always keep in mind that NoSQL database is not built on tables and usually doesn't use structured query language.
- If consistency is mandatory and there will be no drastic changes in terms of the data volume.



# Why MongoDB?

# Why MongoDB



- **It is Open Source :)**
- **High Write Load :**  
MongoDB by default prefers high insert rate over transaction safety. If you need to load tons of data lines with a low business value for each one.
- **High Availability in an Unreliable Environment :**  
Setting replicaSet (set of servers that act as Master-Slaves) is easy and fast. Moreover, recovery from a node (or a data center) failure is instant, safe and automatic.
- **You need to Grow Big:**  
Databases scaling is hard (a single MySQL table performance will degrade when crossing the 5-10GB per table).

# Why MongoDB



- **Location Based:**

MongoDB has built in spacial functions, so finding relevant data from specific locations is fast and accurate.

- **Data Set is Going to be Big (from 1GB) and Schema is Not Stable:**

Adding new columns to RDBMS can lock the entire database in some database, or create a major load and performance degradation in other.


- **You Don't have a DBA:**

If you don't have a DBA, and you don't want to normalize your data and do joins, you should consider MongoDB.

Note: If you are expecting to go big, please notice that you will need to follow some best practices to avoid pitfalls.

<http://www.mongodb.com/use-cases>

# Why MongoDB





# Why MongoDB



## Real World Case Study: Billing



BillRun!


Billing on top of MongoDB

*Ofer Cohen, S.D.O.C. Ltd.*

*MUG IL, February 2014*



mongoDB



# Getting Started

# Getting Started



## Install on Ubuntu:

- Packages MongoDB provides packages of the officially supported MongoDB builds in its own repository. This repository provides the MongoDB distribution in the following packages:
  - mongodb-org
  - mongodb-org-server
  - mongodb-org-mongos
  - mongodb-org-shell
  - mongodb-org-tools
- **Control Scripts** : The mongodb-org package includes various control scripts, including the init script `/etc/init.d/mongod`. These scripts are used to stop, start, and restart daemon processes.
- The package **configures** MongoDB using the `/etc/mongod.conf` file in conjunction with the control scripts.

# Getting Started

---



**Step 1:** Import the public key used by the package management system:

```
sudo apt-key adv --keyserver hkp://keyserver.ubuntu.com:80 --recv 7F0CEB10
```

**Step 2:** Create a list file for MongoDB. using the following command: Create the /etc/apt/sources.list.d/mongodb.list list file


```
echo 'deb http://downloads-distro.mongodb.org/repo/ubuntu-upstart dist 10gen' | sudo tee  
/etc/apt/sources.list.d/mongodb.list
```

**Step 3:** Reload local package database. Issue the following command to reload the local package database:

```
sudo apt-get update
```


**Step 4:** Install the MongoDB packages:

```
sudo apt-get install -y mongodb-org
```



*If you truly wish to master a skill,  
nothing beats hands-on experience*

**SOOOO :)  
Let's Start !**



# Mongo Basics

# Mongo Basics

---



## Binary JSON (BSON)

MongoDB represents JSON documents in binary-encoded format called BSON behind the scenes. BSON extends the JSON model to provide additional data types and to be efficient for encoding and decoding within different languages.

## JSON

JavaScript Object Notation (JSON) is an open, human and machine-readable standard that facilitates data interchange, and along with XML is the main format for data interchange used on the modern web. JSON supports all the basic data types you'd expect: numbers, strings, and boolean values, as well as arrays and hashes.

# Mongo Basics



```
{
  "_id" : 1,
  "name" : { "first" : "John", "last" : "Backus" },
  "contribs" : [ "Fortran", "ALGOL", "Backus-Naur Form", "FP" ],
  "awards" : [
    {
      "award" : "W.W. McDowell Award",
      "year" : 1967,
      "by" : "IEEE Computer Society"
    },
    { "award" : "Draper Prize",
      "year" : 1993,
      "by" : "National Academy of Engineering"
    }
  ]
}
```



# Mongo Basics



## Connect to a mongod:

From a system prompt, start mongo by issuing the mongo command, as follows:

```
$ mongo
```

By default, mongo looks for a database server listening on port 27017 on the localhost interface. To connect to a server on a different port or interface, use the --port and --host options.

## Select a Database:

After starting the mongo shell your session will use the test database by default. At any time, issue the following operation at the mongo to report the name of the current database:

```
> db
```

# Mongo Basics



From the **mongo** shell, display the list of databases, with the following operation:

```
> show dbs
```

Switch to a new database named mydb, with the following operation:

```
> use mydb
```

Confirm that your session has the mydb database as context, by checking the value of the db object, which returns the name of the current database, as follows:

```
> db
```

# Mongo Basics

---




## Display mongo Help

At any point, you can access help for the mongo shell using the following operation:

**> *help***

Furthermore, you can append the ***.help()*** method to some JavaScript methods, any cursor object, as well as the db and db.collection objects to return additional help information.

# Mongo Basics



## Demo